Evaluation of Classification Models

The Limitations of Accuracy

- Consider a 2-class problem:
  - Number of Class 0 examples = 9990
  - Number of Class 1 examples = 10
- If the model predicts everything to be class 0, accuracy is 9990/10000 = 99.9 %
  - Accuracy is misleading!

Measuring model performance

- Do you always want your model to be accurate?

- What other metrics are there?

Confusion Matrix

- We will focus on the predictive capability of a model
  - Rather than how fast it takes to classify or build models, scalability, etc.
- Confusion Matrix:

<table>
<tr><td></td><td colspan="3">PREDICTED CLASS</td></tr>
<tr><td rowspan="3">ACTUAL CLASS</td><td></td><td>Class=Yes</td><td>Class=No</td></tr>
<tr><td>Class=Yes</td><td>a</td><td>b</td></tr>
<tr><td>Class=No</td><td>c</td><td>d</td></tr>
</table>

a: TP (true positive)

b: FN (false negative)

c: FP (false positive)

d: TN (true negative)

Confusion Matrix

	PREDICTED CLASS		
ACTUAL CLASS		Class=Yes	Class=No
	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} = \frac{a + d}{a + b + c + d}$$

- Most widely-used metric

datasciencedojo
unleash the data scientist in you

Accuracy Alternatives

- Precision
  - $P = \frac{TP}{TP+FP} = \frac{a}{a+c}$
  - Sensitive to false positives
- Recall/Sensitivity
  - $R = \frac{TP}{TP+FN} = \frac{a}{a+b}$
  - Sensitive to false negatives
- F1-score
  - $F1 = \frac{2rp}{r+p} = \frac{2a}{2a+b+c}$
  - Harmonic average of precision and recall

		PREDICTED CLASS	
		Class=Yes	Class=No
ACTUAL CLASS	Class=Yes	a (TP)	b (FN)
	Class=No	c (FP)	d (TN)

- Specificity
  - $S = \frac{TN}{FP+TN} = \frac{d}{c+d}$
  - Useful if negative class more important positive

datasciencedojo
unleash the data scientist in you

Controlling Precision and Recall

- What if probabilities are reported?
- Threshold
  - The probability value which separates positive predictions from negative predictions
  - Adjusts class label metrics

| Pid	Prediction	T=0.5	T=0.25	T=0.75
2	.95	Survived	Survived	Survived
3	.86	Survived	Survived	Survived
5	.02	Dead	Dead	Dead
7	.15	Dead	Dead	Dead
13	.48	Dead	Survived	Dead
14	.35	Dead	Survived	Dead
21	.12	Dead	Dead	Dead
24	.01	Dead	Dead	Dead
34	.74	Survived	Survived	Dead
54	.63	Survived	Survived	Dead

ROC (Receiver Operating Characteristic)

- Developed to analyze noisy signals

- TP on the y-axis vs FP on the x-axis

- Plot points for different threshold values

- Curve represents quality of model *independent* of threshold
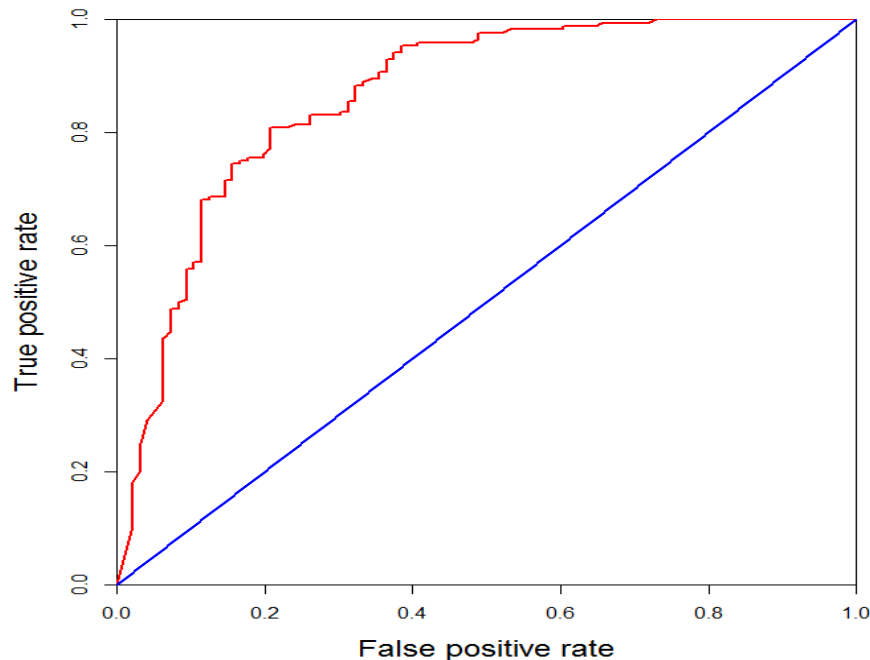
ROC Curve

- Ideal curve (purple)
  - 100% True Positives
  - 0% False Positives
- Random chance (blue)
  - Worst case
- Below diagonal line?
  - Prediction is opposite of the true class

Using ROC for Model Comparison

- No model consistently outperforms the other
  - Purple is better at low thresholds
  - Red is better at high thresholds
- Area Under ROC Curve (AUC)
  - Calculate the area under the curves
  - Compare models directly

WILL MY MODEL BETRAY ME?

datasciencedojo
unleash the data scientist in you

Perils of Overfitting

Data Science Dojo
@DataScienceDojo

Perils of #overfitting @kaggle restaurant revenue prediction Pos 1 drops to 2041 in final ranking.

2041	↑7	Cheng Jiang
2042	↓2041	BAYZ, M.D.
2043	↓81	Alberto

datasciencedojo
unleash the data scientist in you

Overfitting

- The gravest and most common sin of machine learning

- Overfitting: learning so much from your data that you memorize it.
  - You do well on training data
  - But don't do well (or even fail miserably) on test data

Bias/Variance Tradeoff

- You can beat your data to confess anything

Bias/Variance Tradeoff

Methods of Evaluation

- Holdout Set
- Cross validation
- Random subsampling
  - "Repeated holdout"
- Stratified sampling
- Bootstrap
  - Sampling with replacement
  - Discuss later (ensemble)

- We've discussed holdout sets, but not the rest.
- Not limited to just one, apply multiple

Methods: Cross validation

- Most powerful tool for evaluation
- Split dataset into random partitions
  - Stratified sample if appropriate

1	2	3	4	5	6	7	8	9	10

datasciencedojo
unleash the data scientist in you

Methods: Cross validation

- Train model on 2-10, test on 1
- Train (new) model on 1,3-10, test on 2
- Repeat 10 times

1	2	3	4	5	6	7	8	9	10

1	2	3	4	5	6	7	8	9	10

⋮

Methods: Cross validation

- Result: 10 models, labeled by test partition
- Measure bias and variance
- Detect overfitting

	1	2	3	4	5	6	7	8	9	10	Avg	Std
Accuracy	.84	.86	.83	.85	.79	.84	.86	.85	.89	.83	.844	.026
Precision	.79	.78	.81	.79	.85	.76	.82	.71	.75	.76	.782	.040
Recall	.75	.83	.76	.83	.65	.80	.74	.76	.77	.79	.768	.052

Methods: Stratified Sampling

- Used with cross validation or holdout set
- Ensures that all partitions have fixed ratio of classes
  - Same ratio as training set
  - If training set is 5% class 1, 95% class 2, so is each partition
- Use with very uneven class distributions
- Avoid when class distribution isn't constant

QUESTIONS