


# NoSQL Databases

Data Science Dojo

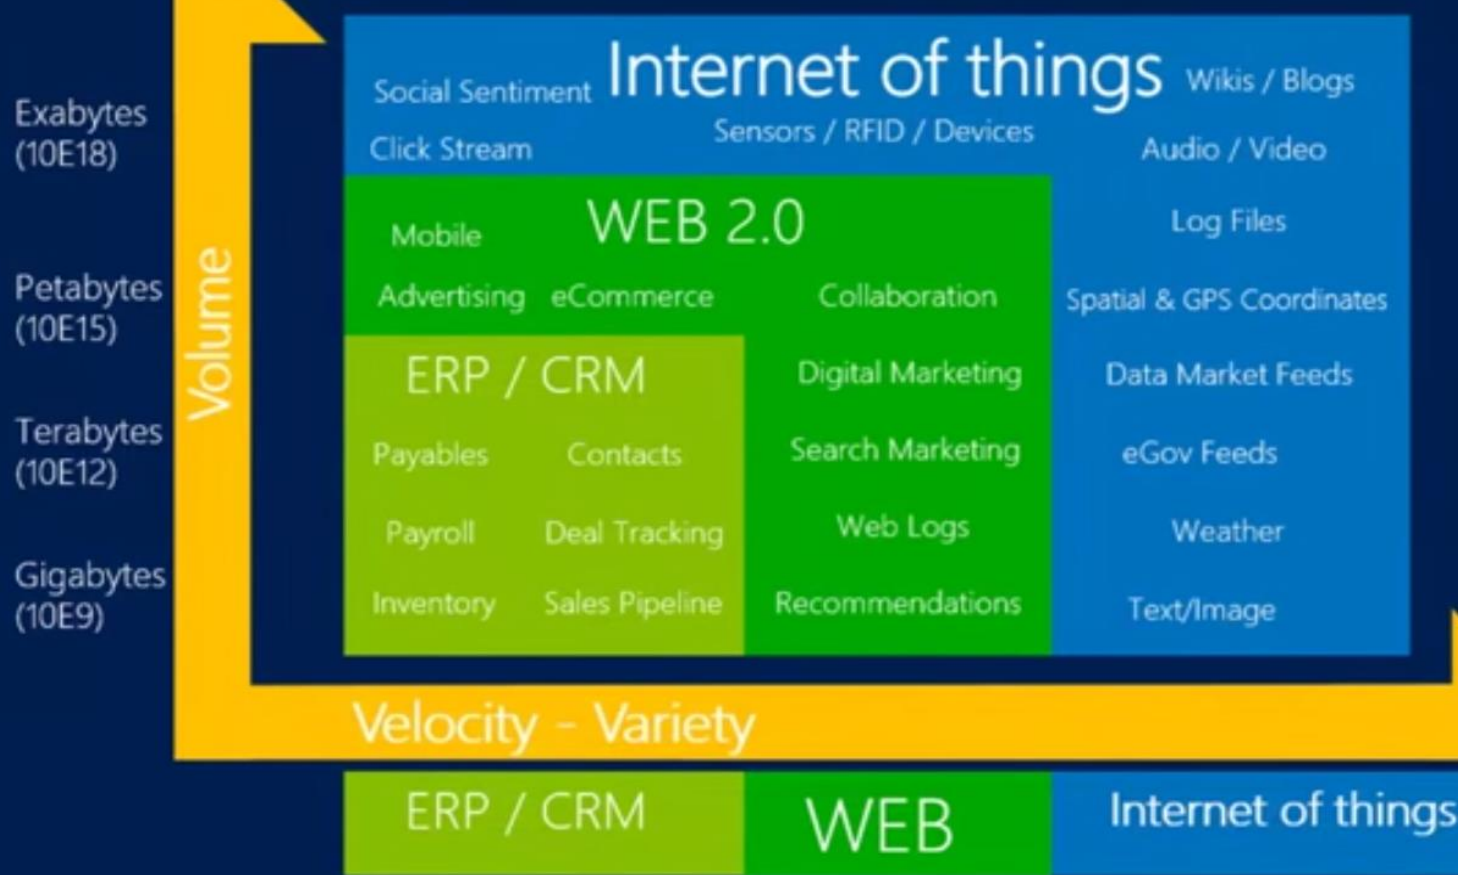
# Scaling, Traditional Relational DB

- Vertical Scaling, scaling UP




# Introducing Big Data

## Continued

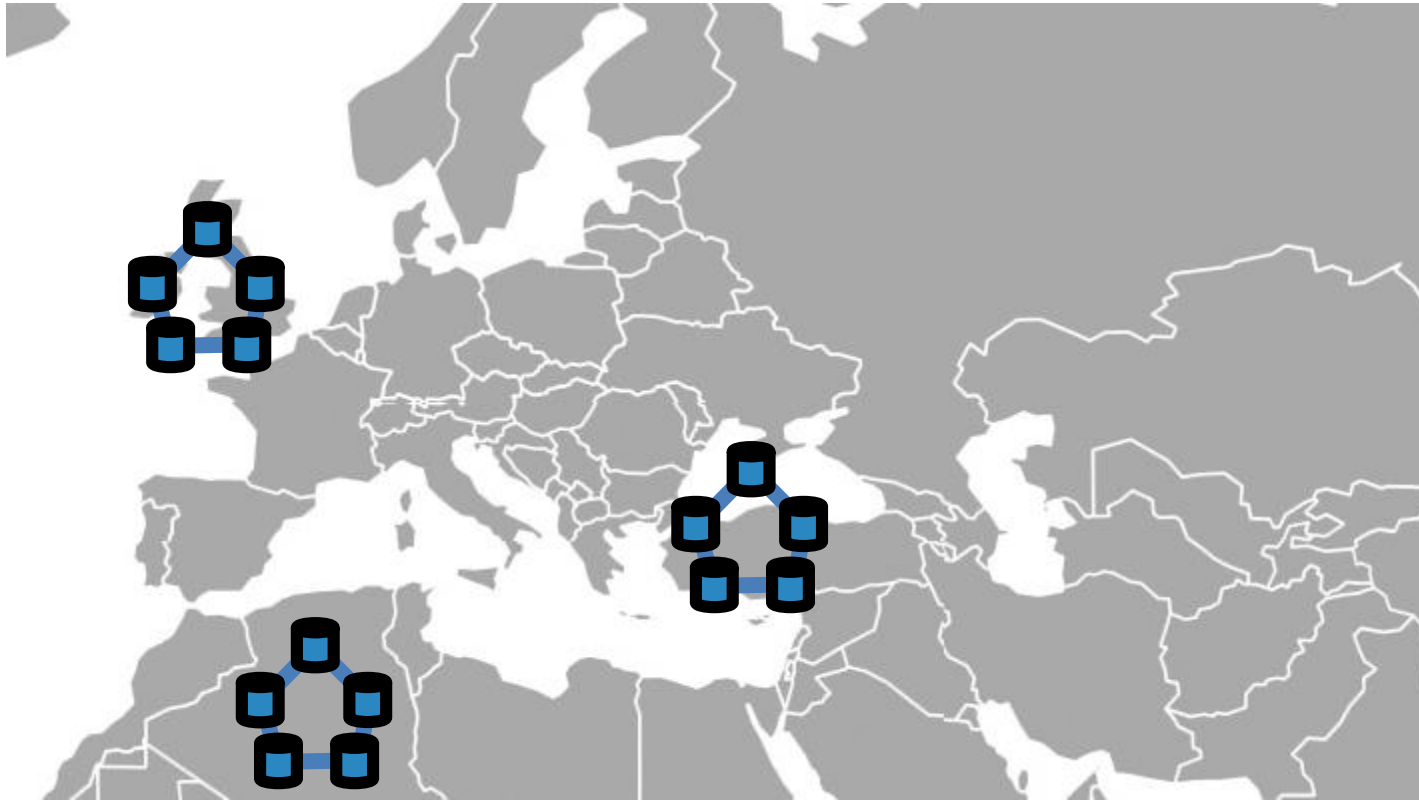


# Scaling, NoSQL Era

- Horizontal Scaling, scaling OUT
- Distributed Databases




# Scaling, NoSQL Era



# Data Architecture

- No standard solution that fits all
- Business and data defines the architecture
- Multiple databases, different types depending on the characteristics of each data subset

# CAP




# CAP Theorem

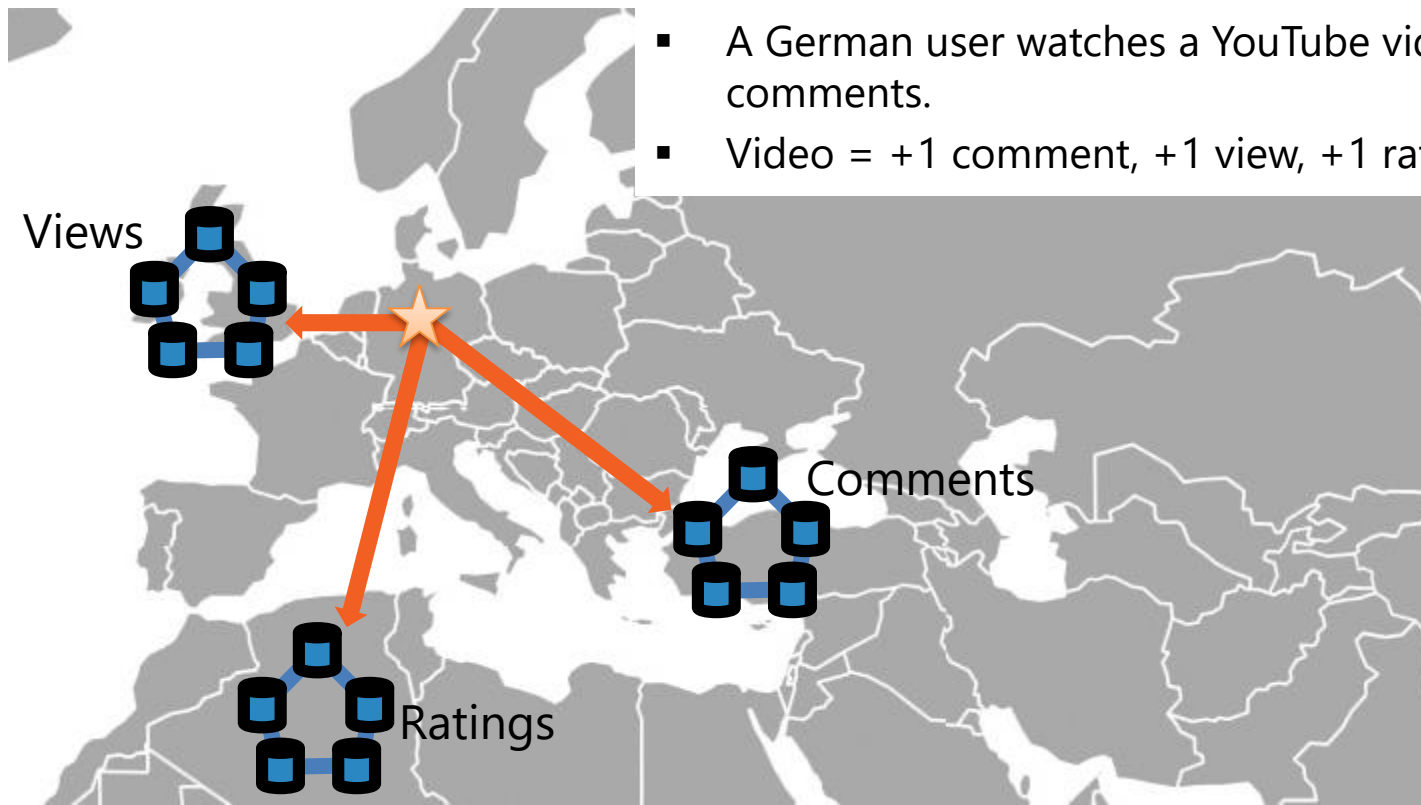
- It is impossible for a distributed processing system to simultaneously provide all three of the following guarantees
  - **Consistency** - A read is guaranteed to return the most recent write for a given client.
  - **Availability** - A non-failing node will return a reasonable response within a reasonable amount of time (no error or timeout).
  - **Partition Tolerance** - The system will continue to function when network partitions occur.



# CAP: Network Partition



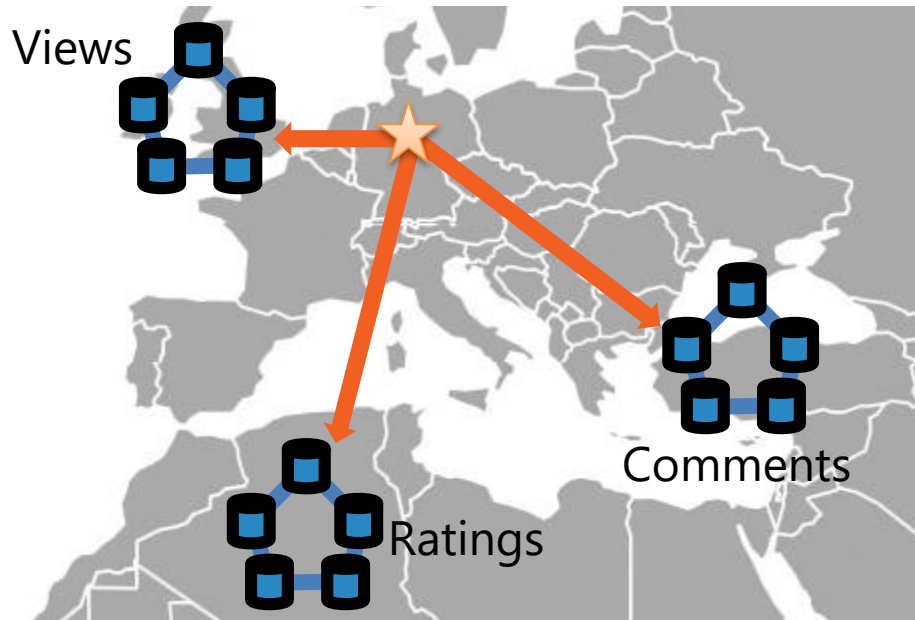
# Wide-area Network Partition



- A German user watches a YouTube video, rates it, then comments.
- Video = +1 comment, +1 view, +1 rating

# CAP: Partition Tolerance

- A German user watches a YouTube video, rates it, then comments.
- Video = +1 comment, +1 view, +1 rating

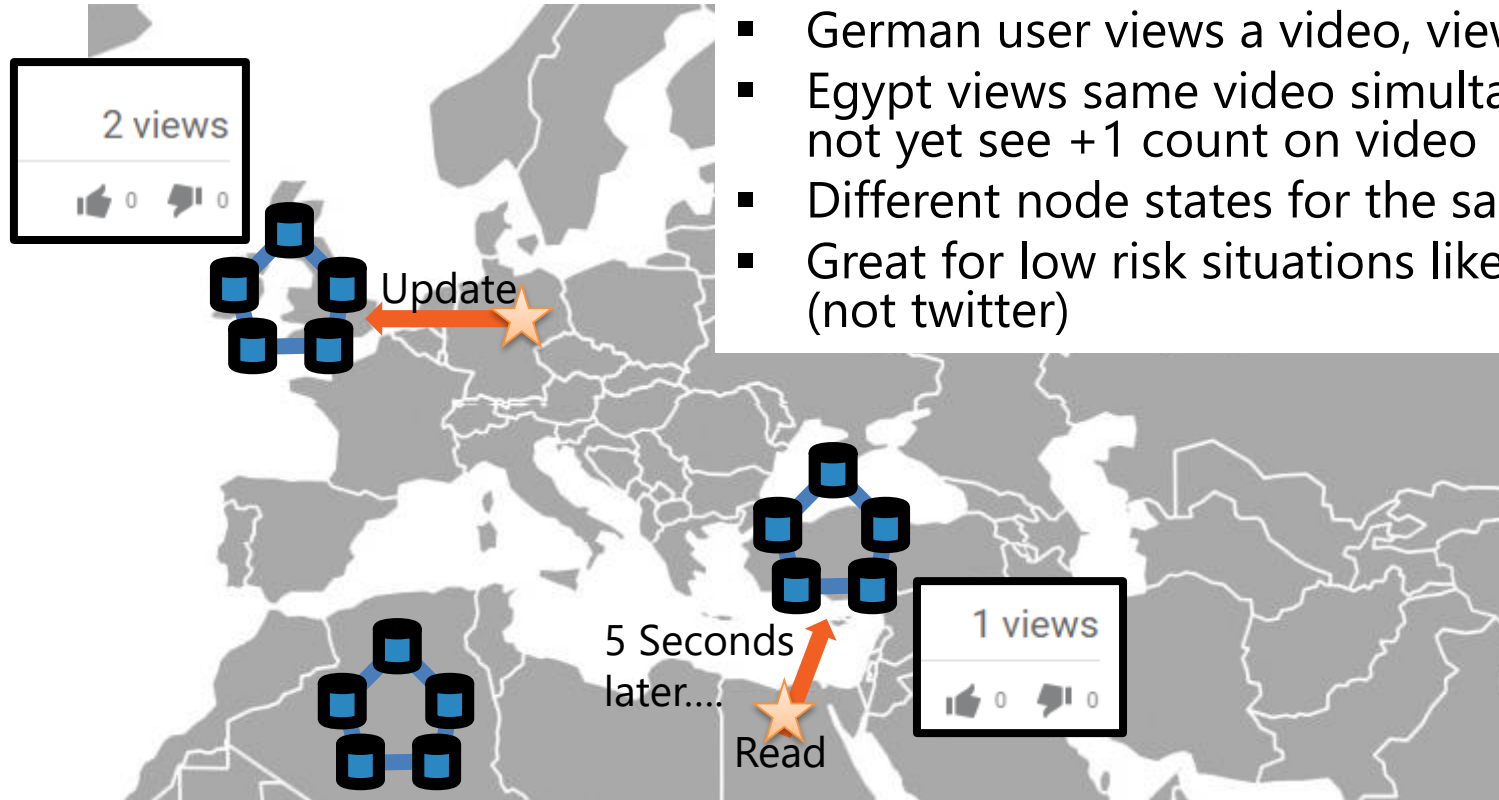


- Definition: The system will continue to function when network partitions occur.
- Most important desirable property for wide-area databases (across geographies)
- Rarest of the 3 desirable properties
  - Must have a distributed and partitioned database

# CAP Theorem

- It is impossible for a distributed processing system to simultaneously provide all three of the following guarantees
  - **Consistency** - A read is guaranteed to return the most recent write for a given client.
  - **Availability** - A non-failing node will return a reasonable response within a reasonable amount of time (no error or timeout).
  - **Partition Tolerance** - The system will continue to function when network partitions occur.

# AP – Lack of (Immediate) Consistency

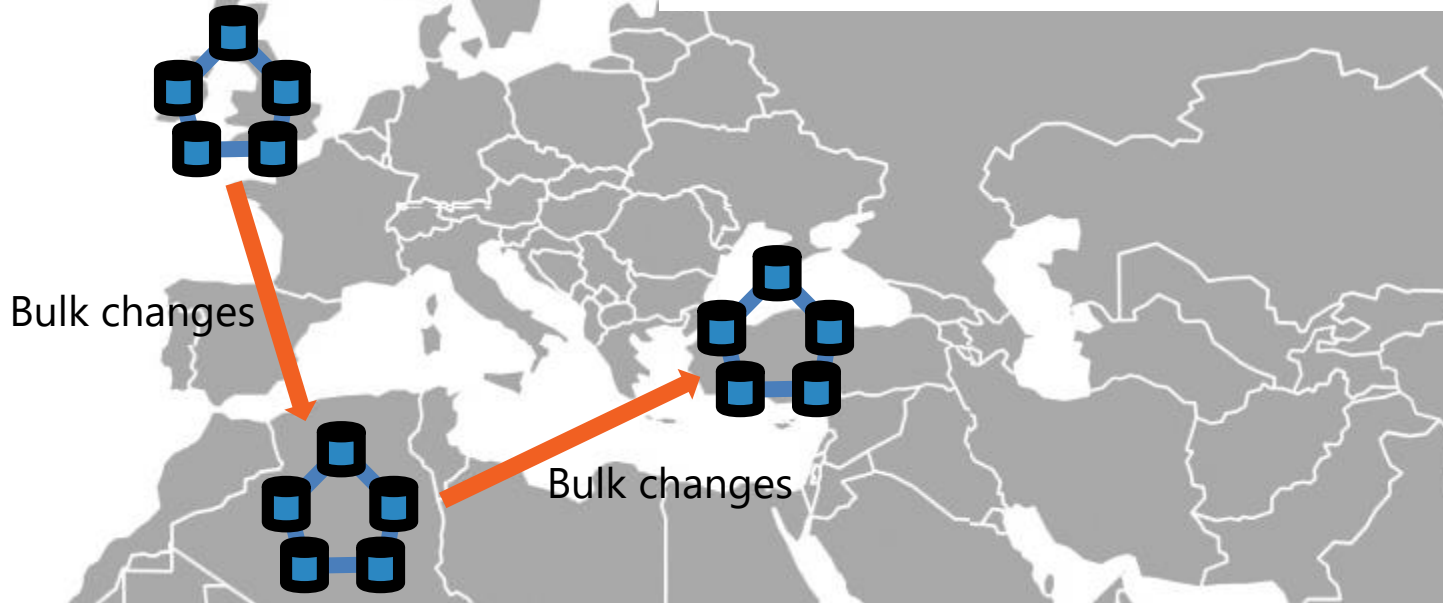


- German user views a video, view count +1
- Egypt views same video simultaneously, does not yet see +1 count on video
- Different node states for the same record
- Great for low risk situations like social media (not twitter)

# AP – Eventual Consistency

1 hour later....

- As time goes by... changes made by German viewer is propagated through the other DB clusters.

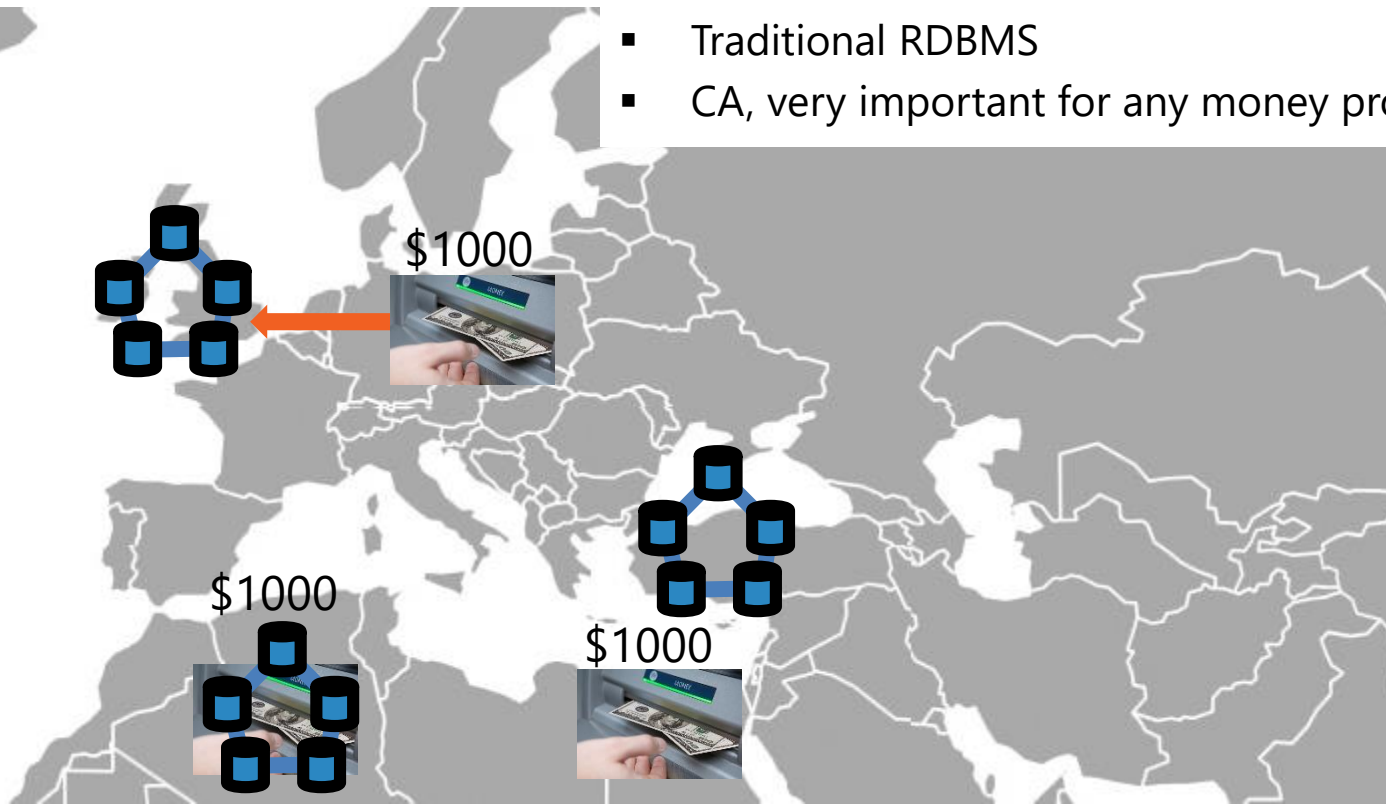


# ATMs



# Why Consistency Matters

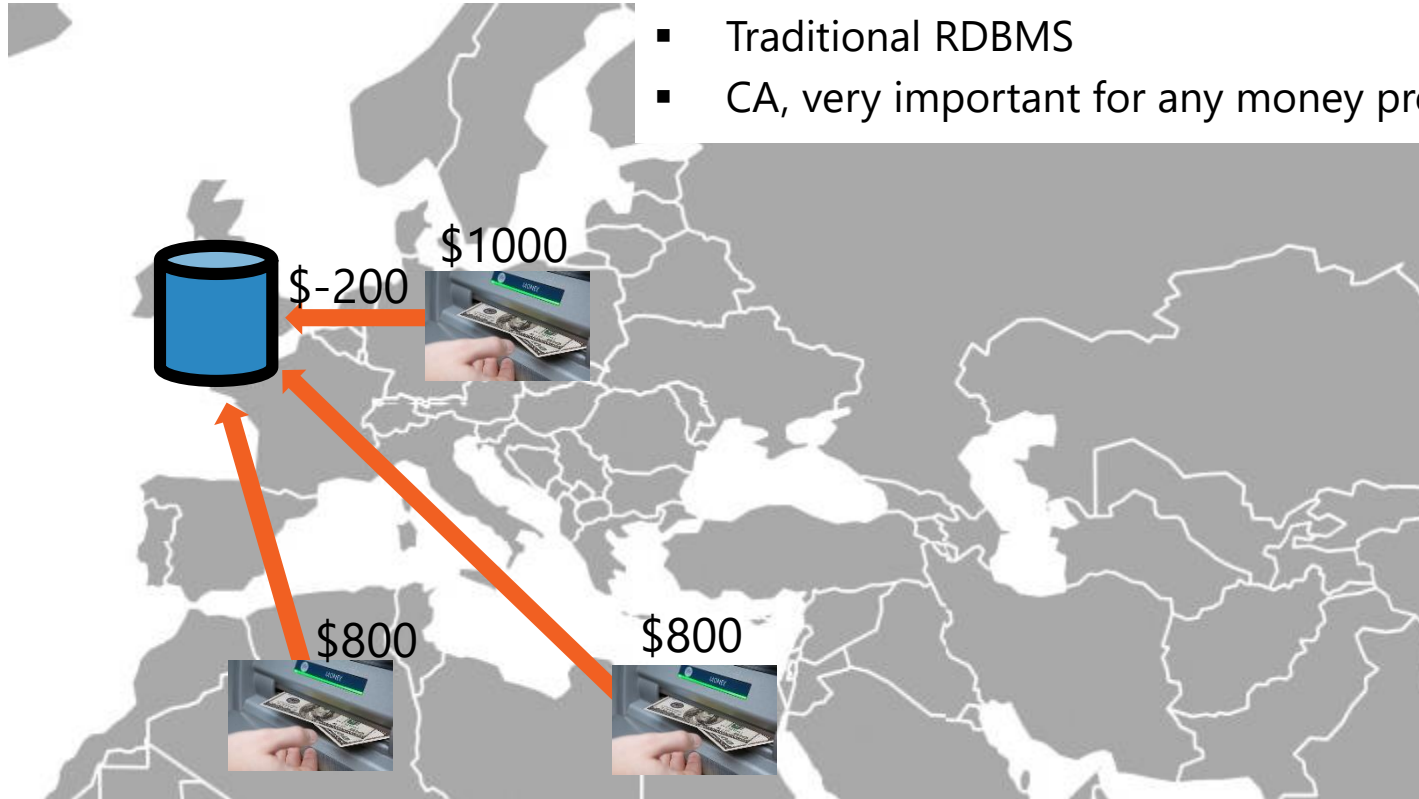
- Traditional RDBMS
- CA, very important for any money processing



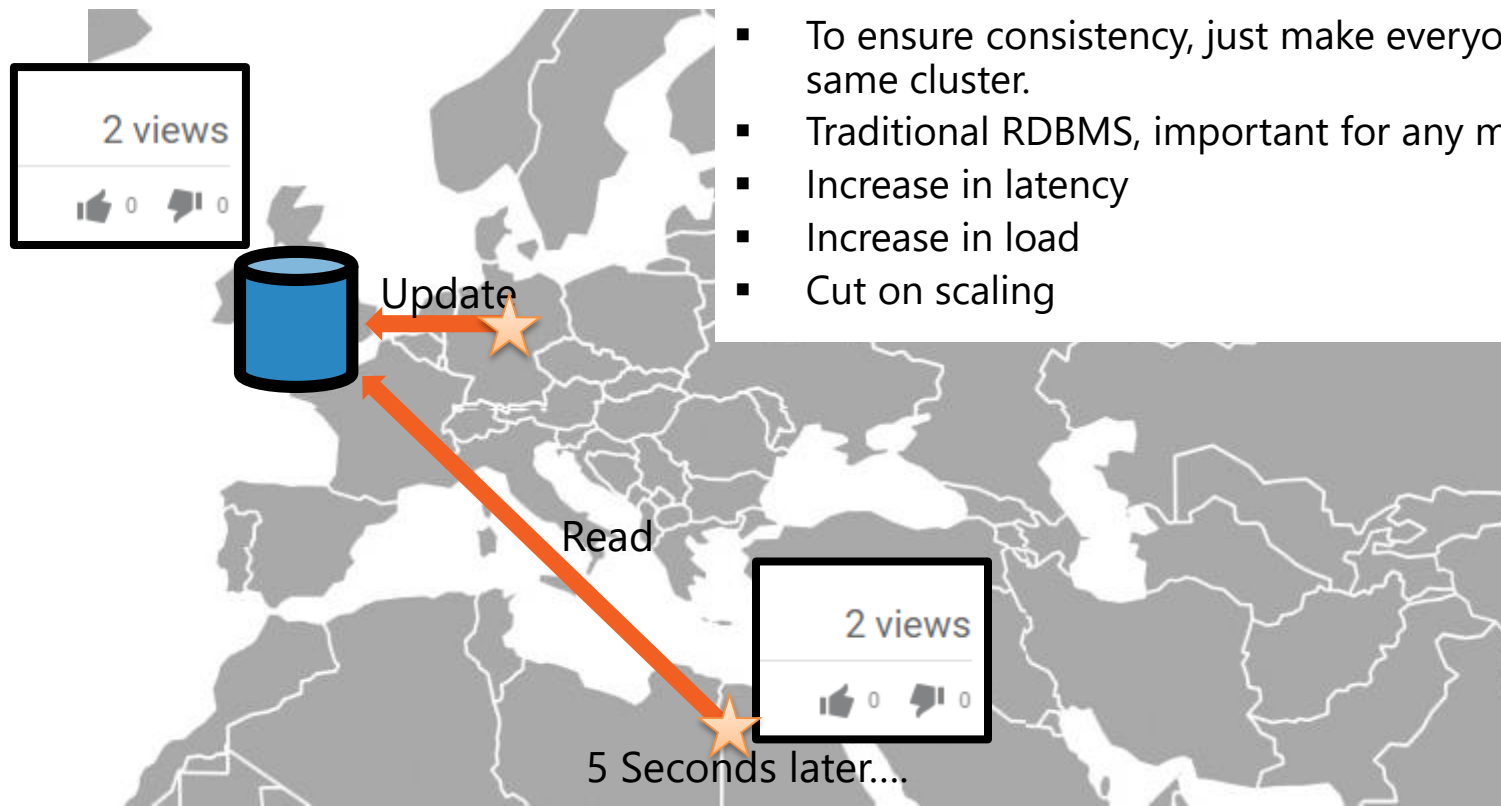


# Why Consistency Matters

- Traditional RDBMS
- CA, very important for any money processing

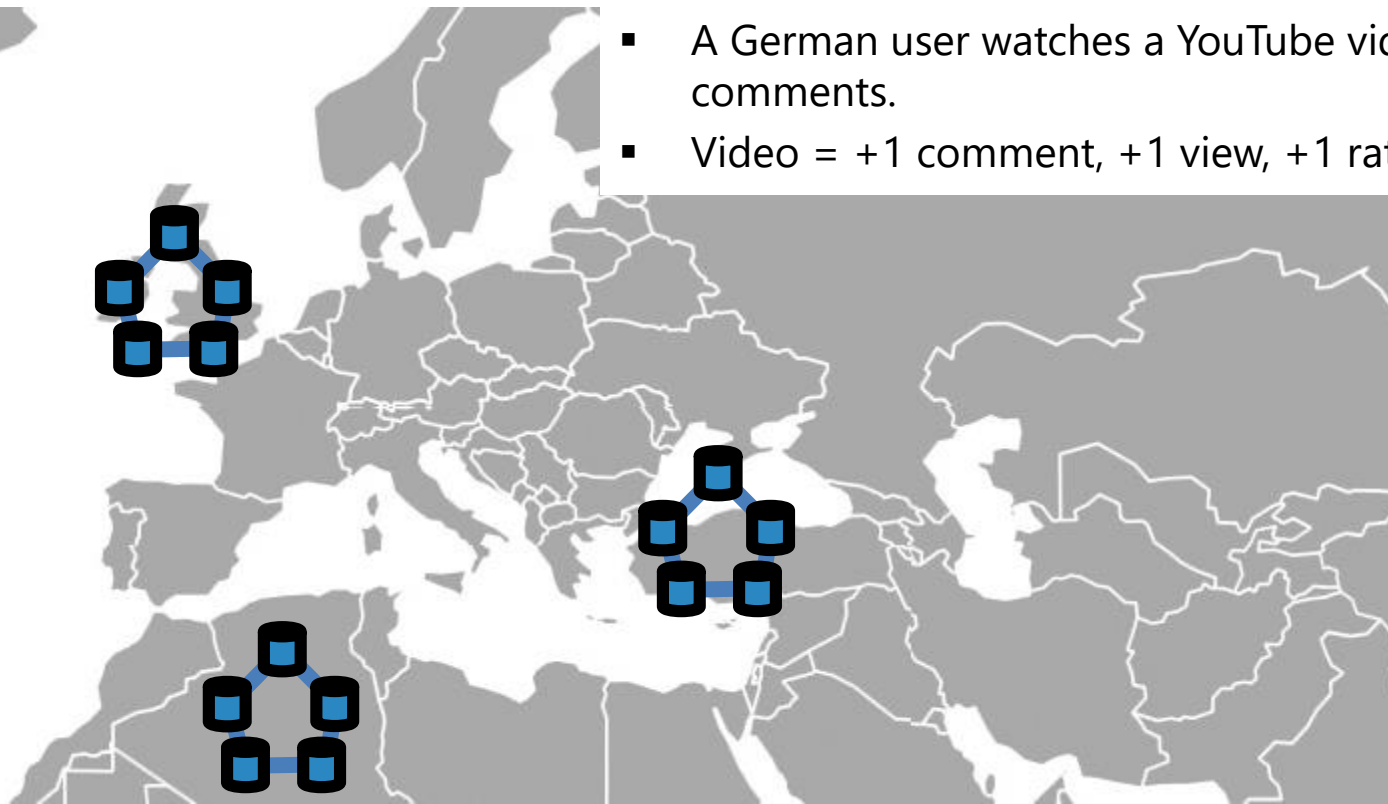


# CA – Lack of Partition Tolerance



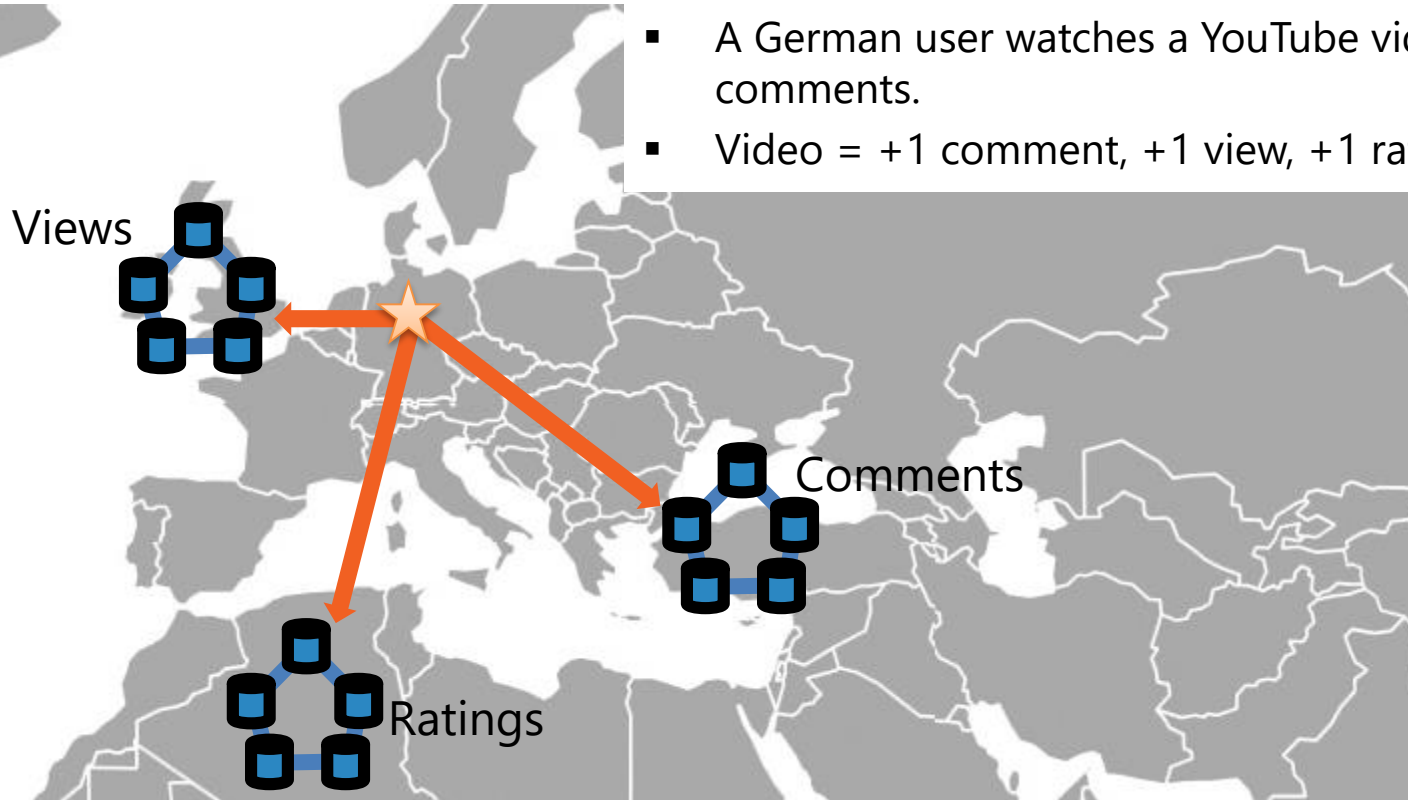
# Consistency + Partitions?

- A German user watches a YouTube video, rates it, then comments.
- Video = +1 comment, +1 view, +1 rating

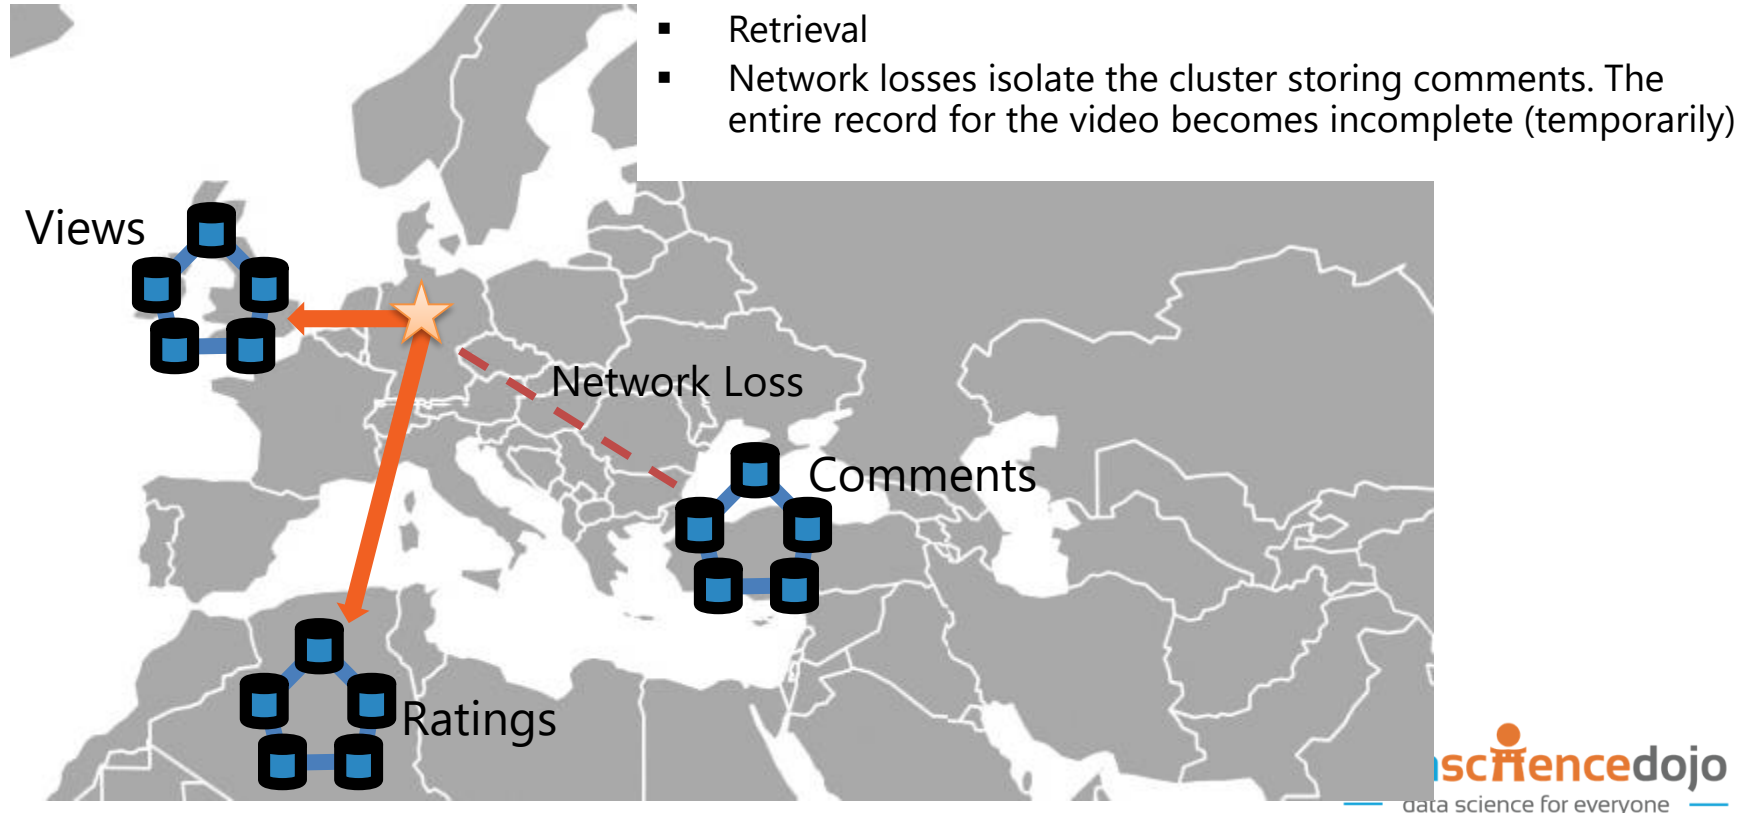


# Consistency + Partitions?

- A German user watches a YouTube video, rates it, then comments.
- Video = +1 comment, +1 view, +1 rating



# CP: Loss of Availability

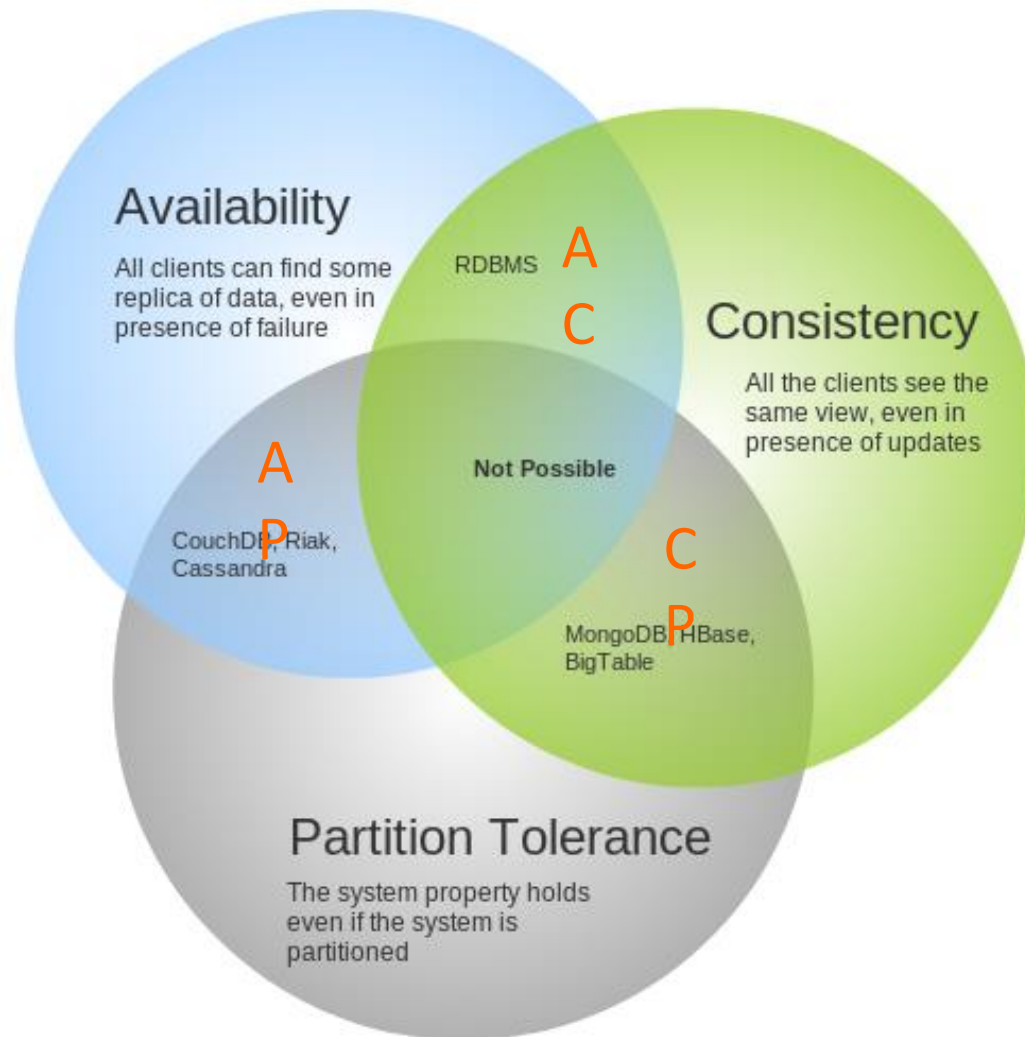


# CAP Theorem

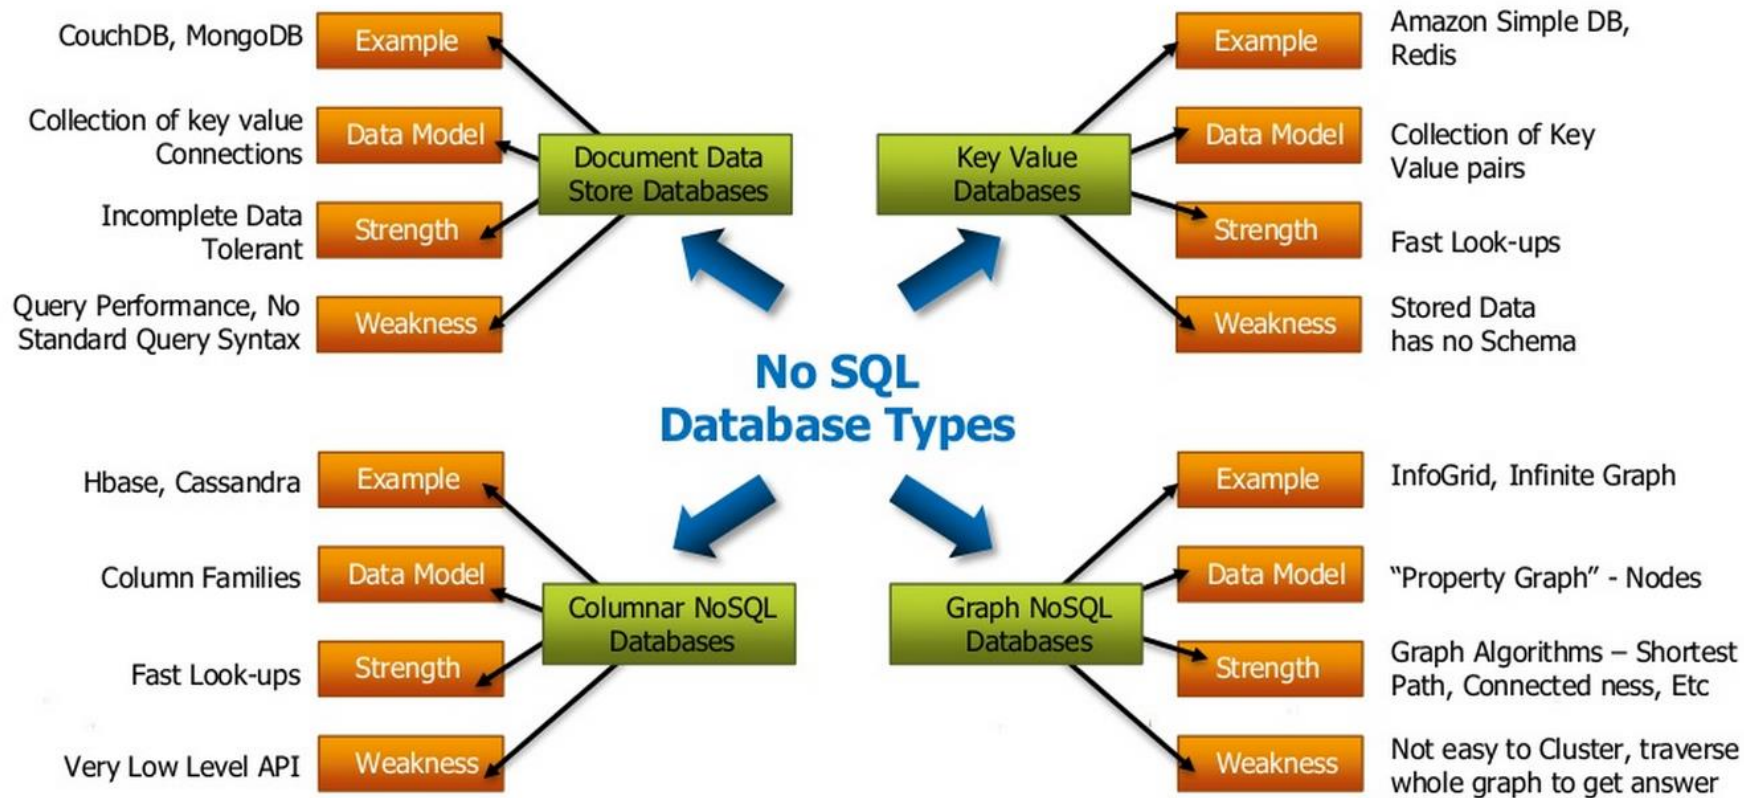
- CAP provides the basic requirements for a distributed system
- **Impossible** to fulfill all 3
- NoSQL databases each make different choices of which 2 to fulfill

# NoSQL vs. SQL

- NoSQL
  - Availability first (Consistency second)
- SQL (Traditional RDBS databases)
  - Consistency first (Availability second)







# What is HBase

- Distributed, non-relational database
  - Columnar, schema-free data model
  - NoSQL on top of Hadoop
- Large scale
  - Linear scalability
  - Billions of rows X millions of columns
  - Many deployments with 1000+ nodes, PBs of data
- Low latency
  - Real-time random read/writes
- Open Source
  - Modeled after Google's BigTable
  - Started in 2006



FLURRY



Adobe



facebook

YAHOO!



Bloomberg

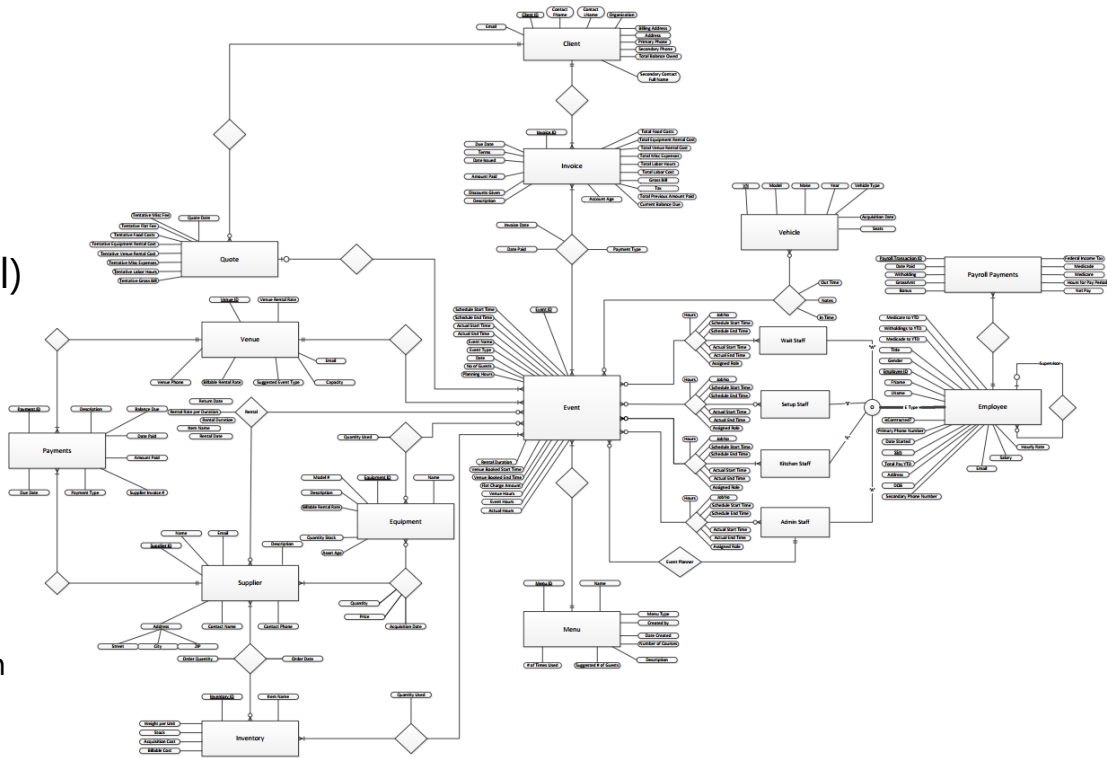
# Problems with Schema

## Pros:

- Relationships
- Consistency: Reading is easier to develop for (each column is guaranteed to be there, even null)

## Cons:

- Records everything, even null
  - Schema validation against each insert
  - Writes can be rejected
- Rigid by design
  - Business processes change faster than databases can be architected and migrated



# Row Store

**Table**

	Country	Product	Sales
Row 1	India	Chocolate	1000
Row 2	India	Ice-cream	2000
Row 3	Germany	Chocolate	4000
Row 4	US	Noodle	500

Pros:

- Fast record query
- Relationships
- Less redundancy
- Single line insert

Cons:

- Thin tables
- Getting a single column value, retrieves the entire record (even null)
  - Terrible with wide tables
- Aggregations must sift through all columns for each row

# Columnar

**Table**

	Country	Product	Sales
Row 1	India	Chocolate	1000
Row 2	India	Ice-cream	2000
Row 3	Germany	Chocolate	4000
Row 4	US	Noodle	500

## Pros

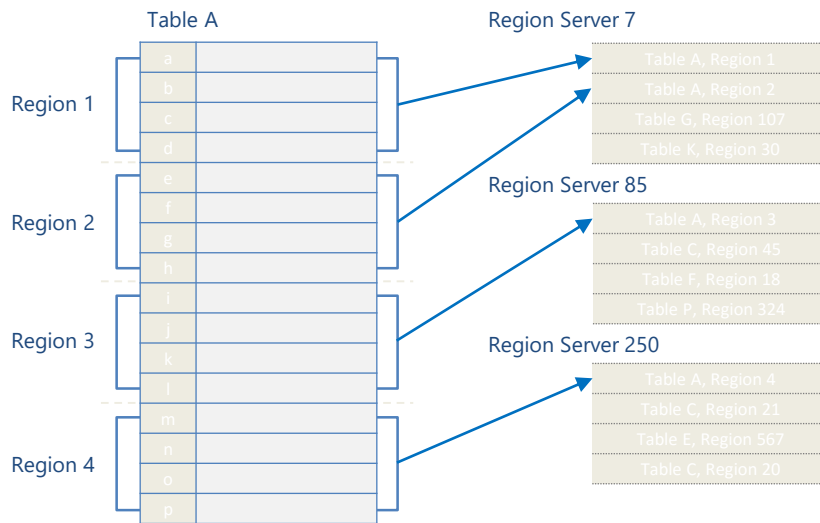
- High speed aggregations
- Compression
- Wide tables are now possible (billions of columns, instead of hundreds)
- High speed snap shot retrieval
- Easily Distributable

## Cons

- Bad for record retrieval
- Terrible at relationships
- O(M) line insert

# Data Model

- Scale-out architecture
  - Automatic sharding of tables
  - Automatic failover
  - Strong consistency for reads and writes
- APIs
  - Get/Put
  - Scan
  - Coprocessors



# Performance Features

- Column Families
- In-memory caching
- High throughput streaming writes

Row Key	Customer		Sales	
Customer Id	Name	City	Product	Amount
101	John White	Los Angeles, CA	Chairs	\$400.00
102	Jane Brown	Atlanta, GA	Lamps	\$200.00
103	Bill Green	Pittsburg, PA	Desk	\$500.00
104	Jack Black	St. Louis, MO	Bed	\$1600.00

Column Families

# Sharding

- Holding rows of database on different partitions
- Same table divided onto different servers, even different geographies
- Reduces index size



# Sharding

- More reliance on interconnection between servers
- Increased latency in querying when more than one shard must be searched
  - Some searches are fast, others are slow
- Often no guarantees about cross shard consistency

# Notable Capabilities

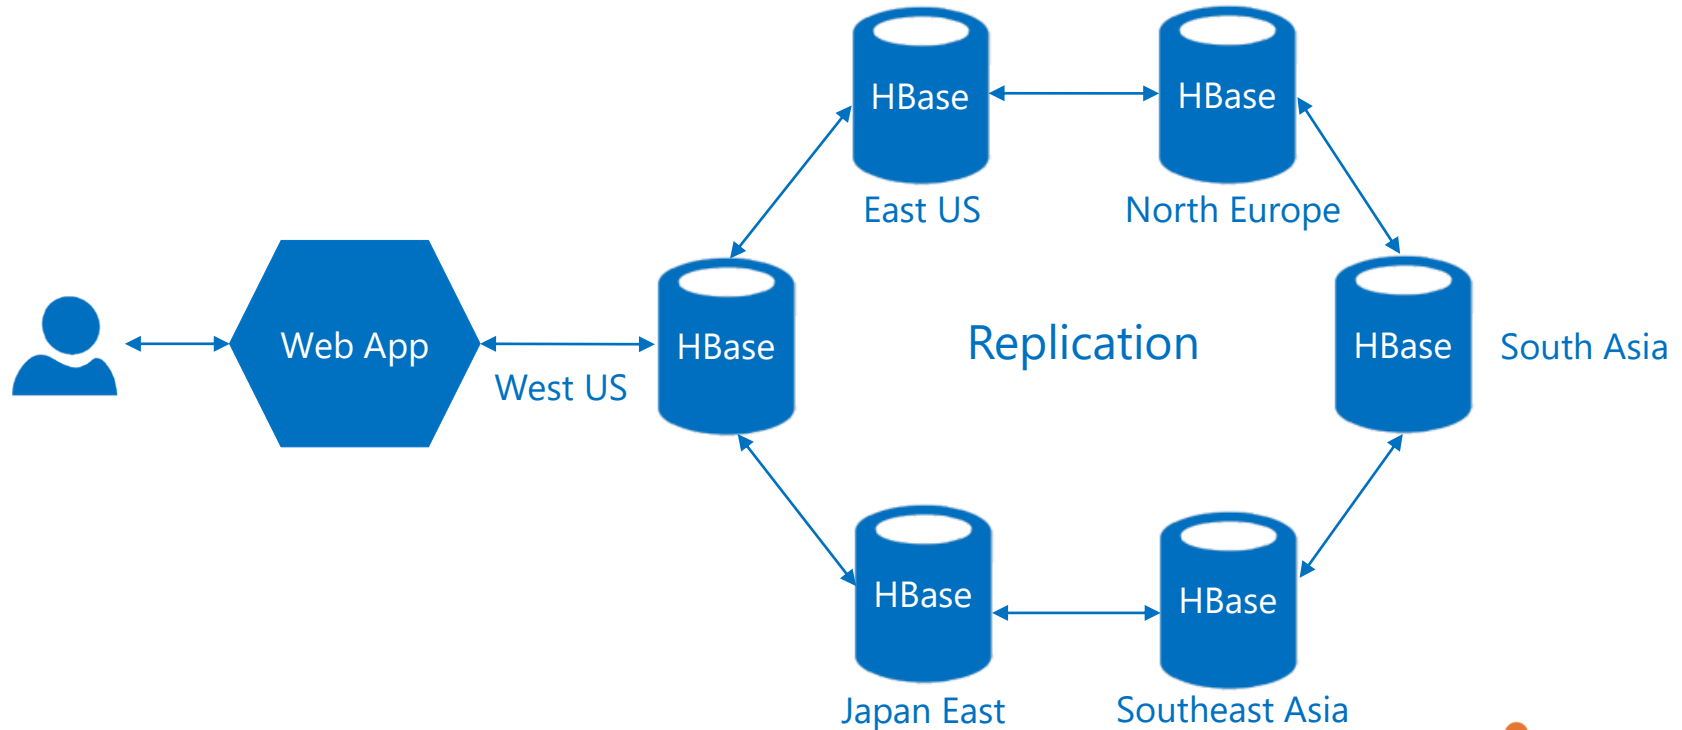
- Integration features
  - Integration with Hadoop MapReduce, Hive, Tez, Spark (hardware pending)
  - Bulk import of large amounts of data
- Client APIs
  - Java, REST, python, node.js, php, .NET

# Use case #1: key value store

- Key value store
  - Message systems
  - Content management systems
- Examples
  - Facebook Messages
  - Twitter-like messages
  - Webtable – web crawler/indexer



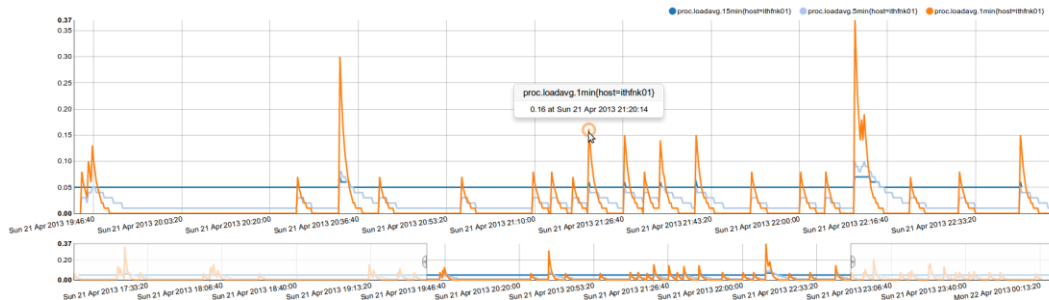
# Use case #1: key value store




- Architecture

# Use case #2: sensor data

- Sensor data
  - Social analytics
  - Time series databases
  - Interactive dashboards with trends, counters, etc
  - Audit log systems



# Use case #2: sensor data



- Architecture

# Questions?